


Primary SFT data composition

1,069 cleaned trajectories; Qwen3 tokenizer; 7.58M context tokens

Examples versus token mass



Sequence length by benchmark

