

Five-epoch training dynamics

Full microbatch logs; loss uses a 25-microbatch rolling mean

